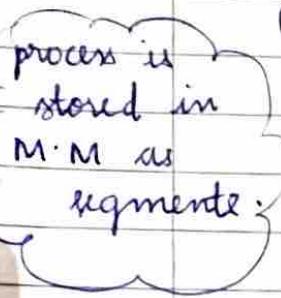


UNIT - 2 Process Management

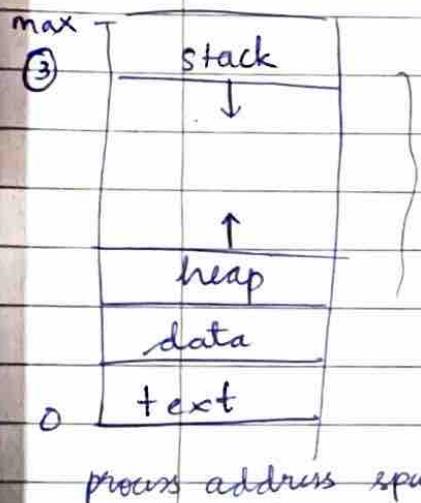
① Memory → Paging (Demand paging)

Date : _____

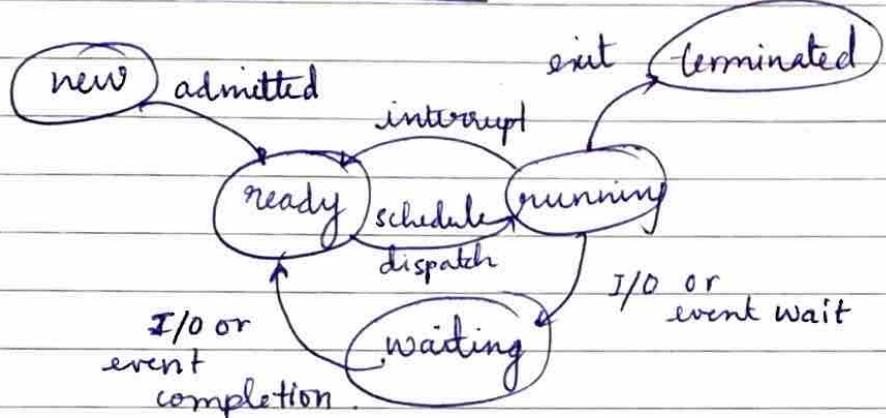
Segmentation - during compile time, divided into
② segments, loaded in M·M segment wise.



- ↳ text seg. - instructions / program code / executable (not declare / initialization) local var.
- ↳ stack seg. - function call returns address parameters
- ↳ data seg. - global variables initialized
- ↳ heap seg. - dynamic memory allocation uninitialized



④ Process State (5 state diagram)



* * *

scheduler
dispatches ready
to running

* once terminated,
removed from
M·M

- ⑤
- new → process is newly loaded into M·M.
 - ready → process that got all the resources required.
 - running → process is being executed.
 - terminated → process completes execution.
 - waiting → process is interrupted while running or event wait or requires non available resource.

⑥

→ moved to
SWAP area (virtual/memory)
secondary

* dynamic binding → when memory is available / free, load and start execution

* static binding → goes to fixed address space

note: cache affinity? no prob problem, solution, merits

* medium
~~middle~~ term scheduler used to select partially executed processes.

* Running → waiting: interrupt, unavailable resource, higher priority process, wait events

Process Control Block (PCB) → Windows

* complete information of all processes in the system.

* each process has one PCB

* also Task Control Block

* stored in secondary memory.

state → any 5 state.

counter number → process ID

PC → address of next instruction

memory limits → memory requirement

e.g.: 1000 to 2000

process state

process number

program counter

registers

memory limits

list of open files

:

open files → f stdio.h, packages.

scheduling info, accounting info, I/O status.
(clock time)

Q) When p1 is executing, & what will be OS be doing? as in which state? (system has one CPU)
one ALU

CPV switching

* when process switches from running to idle, the PCB needs to be stored in memory. (same state)

* When switch back to running,
(reload state)

Process Scheduling Scheduler Types

short term

(CPU scheduling)

→ frequently (ms)

→ allocate CPU

long term

(job scheduling)

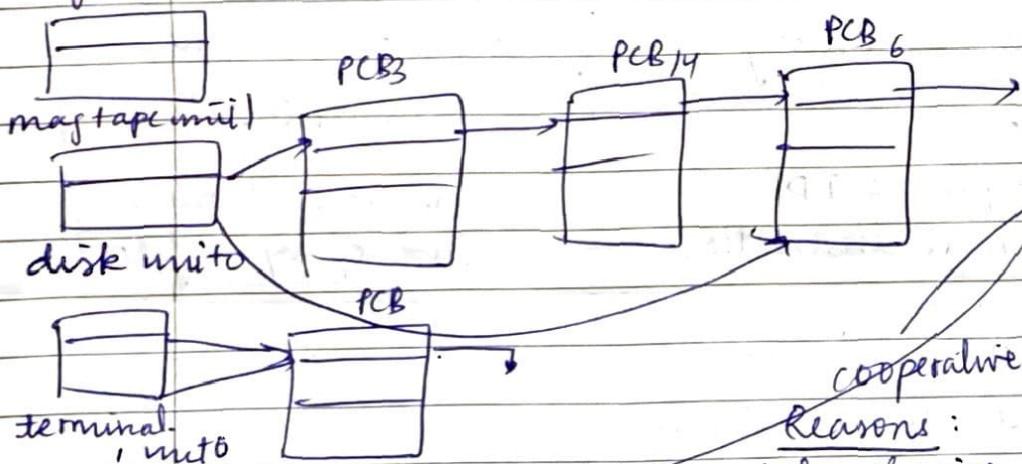
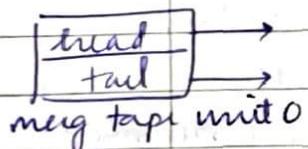
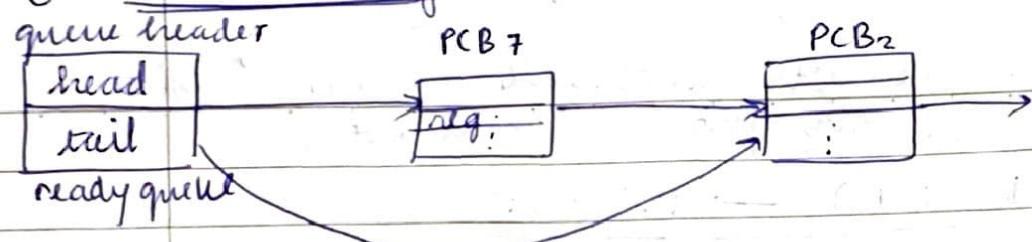
→ infrequently
(s, min)

→ from ready queue

→ controls
multiprogramming

medium term

⑦ ⑧ Process Scheduling



Note:

Date _____
Module is used instead of traditional linear code because M:M is small

Interprocess communication

process

independent dependent

cooperative

Reasons:

- * info sharing
- * computation speed
- * modularity
- * convenience

Two models of IPC:

- Shared memory
- Message passing

Producer consumer buffer.

Bounded

and unbouded
fixed length.

* Initially buffer empty, producer starts.

no acknowledgement
infinite capacity

buffer

shared b/w
producer
and consumer

Message passing

If process P and Q want to commⁿ,

- establish a commⁿ link
- exchange message in send/receive

In PPT, only till buffering

child +
parent --

CPU Scheduling

30/10/24

Date: _____

- * OS needs to utilize CPU to max.
- * CPU is capable of arithmetic and logical
- * Identify how many CPU bound I/O bound instructions in ready queue.
 - eg: CPU I/O + CPU burst time
arithmetic and 80% 20% is time req.
logical.
- * completion time = CPU burst + I/O burst time
- * Short-term scheduler is used to the algorithms we are learning.
- * CPU scheduling takes place when
 - ① running → waiting
 - ② waiting to ready
 - ③ running → ready
 - ④ terminates
- * Non-preemptive - can't continue so scheduler has to be called.
- * Preemptive - (2) & (3)

Scheduling criteria

- * CPU utilization → uses to max extend
- * Turnaround time → no of instructions complete per time
- * Waiting time → process waiting in ready queue.
- * Response time → time b/w job submission & completion
- (first response it got when selected)

1) First-Come, First-Served (FCFS) Scheduling (non-preemptive)

Process	Burst time (ms)	eq: order: P ₁ , P ₂ , P ₃
P ₁	24	Step 1: Gantt chart
P ₂	3	P ₁ P ₂ P ₃
P ₃	3	24 27 30

Arrival time → enter the ready queue (one by default)

Step 2: Arrival time	Waiting time		Turnaround time
	P ₁	P ₂	P ₃
	0	24	24
	24	27	27
	27	30	30
avg	17	27	27

eg:

	P1	8	Step 1:	P1	P2	P3	P4	Date: 21 26
P1	8			8	12	21	26	
P2	4							
P3	9							
P4	5							

note:

$$\boxed{TAT = \text{waiting time} + \text{burst time}}$$

$$\boxed{\text{Waiting time} = \text{start time} - \text{arrival time}}$$

Step 2: Waiting time Step 3: Turnaround time (TAT)

P1	0	8
P2	8	12
P3	12	21
P4	21	26
avg	10.25	16.75

eg:

	P1	8	0	Waiting time	TAT
P1	8	4	1	$P1 = 0 - 0 = 0$	$P1 = 0 + 8 = 8$
P2	4	9	2	$P2 = 8 - 1 = 7$	$P2 = 7 + 4 = 11$
P3	9	5	3	$P3 = 12 - 2 = 10$	$P3 = 10 + 9 = 19$
P4	5	1	3	$P4 = 21 - 3 = 18$	$P4 = 18 + 5 = 23$
avg				$\underline{\underline{= 8.75}}$	$\underline{\underline{15.25}}$

- * easy to code
- * not effective process

[2] shortest Job First Scheduling (non preemptive)

- eg:
- * gives minimum waiting time
 - * choose job with least burst time

① order: P2, P4, P1, P3

	P1	8	P2	P4	P1	P3
P1	8		6	4	9	17
P2	4					
P3	9					
P4	5					

Step 2: Waiting time TAT

P1	=	9	17
P2	=	0	4
P3	=	17	26
P4	=	4	9
avg		$\underline{\underline{7.5}}$	$\underline{\underline{14}}$

eg:

	P1	24	P2	P3	R1	Date: 30 30
P1	3		0	3	6	
P2		3			6	
P3						
			wt		TAT	
P1			6		30	
P2			0		3	
P3			3		6	
avg			3		13	

4/11/24

3) Shortest Job First (preemptive)

L, Shortest Remaining Time First Algorithm.

eg: ① Ready queue

	CPU Burst time	Arrival time	When new process, scheduling scheduler takes a break, and decides whether to continue or choose new process (preempt old old process)
P1	8	0	
P2	4	1	
P3	9	2	
P4	5	3	

P1	P2	P4	P1	P3
0	1	5	8 10	17 25

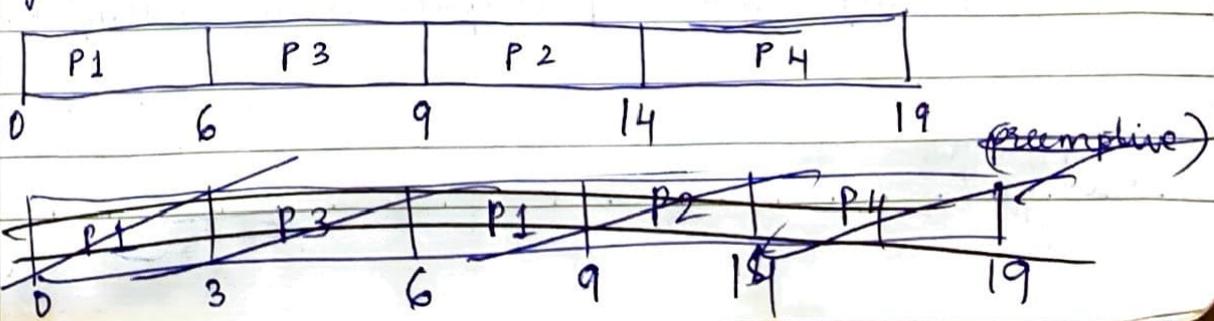
	Waiting time	TAT
P1	$10 - 1 - 0 = 9$	$9 + 8 = 17$
P2	$1 - 1 - 0 = 0$	$0 + 4 = 4$
P3	$17 - 2 = 15$	$15 + 9 = 24$
P4	$5 - 3 = 2$	$2 + 5 = 7$

avg 6.5 @ 13

②

	Arrival time	CPU burst time
P1	0.0	6 3
P2	3.0	5 3
P3	3.0	3 1
P4	5.0	5 5

Gantt chart



will be same for preemptive and non-preemptive (non-preemptive)

	<u>Waiting Time</u>	<u>TAT</u>	Date:
P1	0	$0 + 6 = 6$	
P2	$9 - 3 = 6$	$6 + 5 = 11$	
P3	$6 - 3 = 3$	$3 + 3 = 6$	
P4	$14 - 5 = 9$	$9 + 5 = 14$	
avg.	4.5	9.25	

Note: Non preemptive SJF if when arrival time is diff

P1	8	0
P2	4	1
P3	9	2
P4	5	3

P1		P2		P4		P3		
0		8		12		17		26

	<u>Waiting</u>	<u>TAT</u>
P1	$0 - 0 = 0$	$0 + 8 = 8$
P2	$8 - 1 = 7$	$7 + 4 = 11$
P3	$17 - 2 = 15$	$15 + 9 = 24$
P4	$12 - 3 = 9$	$9 + 5 = 14$
avg:	7.75	14.25

- [4] Priority Scheduling algorithm
 * CPU is allocated to highest priority process from ready queue. Each process has priority no. and it is unique.
 ↳ sometimes if 2 or more processes have same priority. Then FCFS algo. is applied for solving the tie.

- * In case studies, we assume that lowest no. of have highest priority.
- * Priority can be preemptive or non preemptive.

Multiprocessor Scheduling

eg: Non preemptive

RQ	CU BT	Priority
P1	3	2
P2	6	4
P3	4	1
P4	2	3

Arrival time of all = 0.

	P3	P1	P4	P2	Date:
0	4	7	9	15	
		W.T			TAT
P1	4				$4+3=7$
P2	9				$9+6=15$
P3	0				$0+4=4$
P4	7				$7+2=9$
avg	5				8.75

If arrival time is 0, 1, 2, 3

	P1	P3	P4	P2			
	0	3	7	6	15	WT	TAT
P1	0				P1	$0-0=0$	$0+3=3$
P2	$6-1=5$				P2	$9-1=8$	$8+6=14$
P3	$3-2=1$				P3	$3-2=1$	$4+4=8$
P4	$9-3=6$				P4	$7-3=4$	$4+2=6$
avg	1.75				avg	3.25	7

note: For priority, focus on Priority column only

eg: Preemptive

	BT	Priority	AT
P1	10	3	0
P2	1	1	1
P3	2	3	2
P4	1	4	3
P5	5	2	4

Gantt chart.

since P1 is in swap
and will take extra clock
cycle to become ready,
P3 is chosen

this is only
when they
have same
priority

	W.T	T.A.T
P1	$9-0-1=8$	$8+10=18$
P2	$1-1=0$	$0+1=1$
P3	$2-2=0$	$0+2=2$
P4	$18-3=15$	$15+1=16$
P5	$4-4=0$	$0+5=5$
avg	4.6	8.4

process	B.T	Priority	A.T	① Non preemptive FCFS				
P1	7	3	0	P1	P2	P3	P4	P5
P2	2	2	3	0	1	3	5	
P3	3	1	4			WT		
P4	1	1	4	P1	P2	3		
P5	3	3	5	P3	P4	5		

	WT	TAT
P1	0 - 0 = 0	0 + 7 = 7
P2	7 - 3 = 4	4 + 2 = 6
P3	9 - 4 = 5	5 + 3 = 8
P4	12 - 4 = 8	8 + 1 = 9
P5	13 - 5 = 8	8 + 3 = 11
avg	5	8.2

① Non preemptive FCFS (arrival order)

P1	P2	P3	P4	P5
0	7	9	12	13

	WT	TAT
P1	0 - 0 = 0	0 + 7 = 7
P2	8 - 3 = 5	5 + 2 = 7
P3	10 - 4 = 6	3 + 6 = 9
P4	7 - 4 = 3	3 + 1 = 4
P5	13 - 5 = 8	8 + 3 = 11
avg	4.4	7.6

② Non preemptive SJF (burst-time order)

P1	P2	P3	P4	P5
0	3	5	8	10

	WT	TAT
P1	12 - 0 - 3 = 9	9 + 7 = 16
P2	3 - 3 = 0	0 + 2 = 2
P3	6 - 4 = 2	2 + 3 = 5
P4	5 - 4 = 1	1 + 1 = 2
P5	9 - 5 = 4	4 + 3 = 7
avg	3.2	6.4

③ Preemptive SJF (burst-time order)

P1	P2	P4	P3	P5	P1	
0	3	5	6	9	12	16

	WT	TAT
P1	12 - 3 - 0 = 9	9 + 7 = 16
P2	8 - 3 - 1 = 4	4 + 2 = 6
P3	4 - 4 = 0	0 + 3 = 3
P4	7 - 4 = 3	3 + 1 = 4
P5	9 - 5 = 4	4 + 3 = 7
avg	3.2	6.4

④ Preemptive Priority (priority order)

P1	P2	P3	P4	P2	P5	P1	
0	3	4	7	8	9	12	16

	WT	TAT
P1	12 - 3 - 0 = 9	9 + 7 = 16
P2	8 - 3 - 1 = 4	4 + 2 = 6
P3	4 - 4 = 0	0 + 3 = 3
P4	7 - 4 = 3	3 + 1 = 4
P5	9 - 5 = 4	4 + 3 = 7
avg	4	7.2

⑤ Non preemptive priority (priority order)						W.T	TAT	
P1	P3	P4	P2	P5		P1	0 - 0 = 0	0 + 7 = 7
7	10	11	13	16	P3	11 - 3 = 8	8 + 2 = 10	D _{arr}
					P3	7 - 4 = 3	3 + 3 = 6	
					P4	10 - 4 = 6	6 + 1 = 7	
					P5	13 - 5 = 8	8 + 3 = 11	
					avg	5.2	8.2	8.2

5 Round Robin Scheduling Algorithm

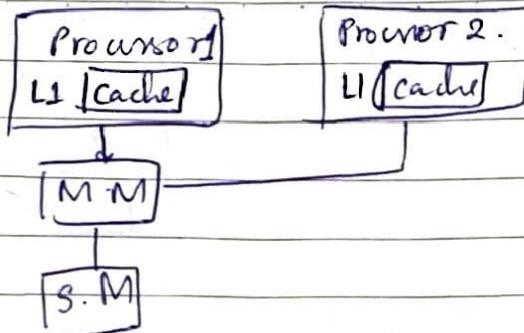
- * shares CPU with all processes equal amount of time
 - * process selection is like FCFS
 - * Preemptive by nature.

Ques:	Process	CPU burst time	Time slice = 2 msec
	P1	6	
	P2	3	
	P3	2	
	P4	5	
			0 2 4 6 8 10 11 13 15 16
			WT TAT
	P1	$13 - 2 - 2 = 9$	$9 + 6 = 15$
	P2	$10 - 2 = 8$	$8 + 3 = 11$
	P3	$4 - 0 = 4$	$4 + 2 = 6$
	P4	$15 - 2 - 2 - 0 = 11$	$11 + 5 = 16$
		avg 8	12

eg:	P1	7	0	Time slice: 1ms
	P2	3		
	P3	4		
	P4	1	4	0 1 2 3 4 5 6 7 8 9 10 11 12 13
	P5	5		[P5 P1 P1] 13 14 15 16
				WT TAT
	P1	15 - 6 = 9	9 + 7 = 16	
	P2	8 - 1 - 3 = 4	4 + 2 = 6	
	P3	12 - 2 - 4 = 6	6 + 3 = 9	
	P4	5 - 4 = 1	1 + 1 = 2	
	P5	13 - 2 - 5 = 6	3 + 6 = 9	
	avg	5.2	8.4	

Process synchronization

- two or more dependent processes are working on same data.



Producer

```

while (true) {
    while (counter == SIZE)
        /* do nothing */
    buffer[i] = next prod.
    i = i + 1;
    in = (in + 1) % SIZE
    counter += 1
}
  
```

Consumer

```

while (true) {
    while (counter == 0)
        /* do nothing */
    next-con = buffer[out]
    out = (out + 1) % SIZE
    counter -= 1
}
  
```

* Race condition: while producer is executing, consumer starts execution causing inconsistency in data.

Solution: Synchronization

* Critical section: Part of M-M reserved for process coordination.

If any var has its value continuously changed, those are stored in critical section.

Solution:

① Mutual exclusion

② Progress

③ Bounded waiting

ask permission

critical section

exit section

remainder

section

{while(true)}

* Peterson's solution:

P1:

P2:

shared variables

int turn;

boolean flag[2];

```

do {
    flag[i] = true;
    turn = i;
    while (flag[j] && j != i)
        turn = i;
    flag[i] = false;
} remainder
  
```

```

do {
    flag[j] = true;
    turn = j;
    while (flag[i] && i != j)
        turn = j;
    flag[j] = false;
} remainder
  
```

* only two processes applicable

{while(true);}

{while(true)}

Synchronization Hardware

* locking - protecting critical section by locks.

Date: 13/11/24

test-and-set instruction

```
boolean test_and_set(boolean *target) {
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

} executed as one
instruction - h/w
solution
CPU instruction

① Executed atomically

② Returns original value of passed parameter

③ Set the new value of passed parameter as TRUE

Q) do { while (test_and_set(&lock));
 // do nothing
 /* critical section */
 lock = false;
 /* remainder section */
} while (true);

lock is shared variable

(Ans) P1 enters critical section
 * Execute test-and-set by passing lock.
 * Initially lock is FALSE
 * make lock = TRUE and return false. so P1 enters critical section.
 * once done, lock again becomes false.

Q) If P1 is in critical section, and P2 wants enter critical section, will P2 be allowed or denied?

No because, * when P2 wants to enter, it executes test-and-set with lock value TRUE.

* makes lock = TRUE and returns true.

* hence P2 nothing happens and P2 doesn't enter critical section.

Q) If P1 is in remainder section, and it wishes to enter critical section, is it allowed or denied?

compare-and-swap instruction

```
int CAS (int* value, int expected, int new-val)
{
    int temp = *value;
    if (*value == expected)
        *value = new-val;
    return temp;
}
```

① execute ~~enters~~ atomically
② Returns old val of value.
③ if val == expected, set val to passed value of new-val.

```
do {
    while (CAS (&lock, 0, 1) != 0);
    // do nothing
    /* critical section */
    lock = 0; /* removes lock */
    // write (some);
}
```

case i) If P1 wants to enter CS.

- * P1 executes CAS with lock initially '0', expected is 0, new-val is 1.
- * make lock as 1 since lock == expected and returns 0.
- * while loop will be 0!=0 so P1 enters CS.
→ fails.
- * once done lock becomes 0.

case ii) P2 wants to enter CS.

- * P2 executes CAS with lock = 1, expected = 0, new-val = 1.
- * Since lock ≠ expected, lock value doesn't change.
- * so CAS returns 1. In loop condition 1 != 0 becomes true, so it does nothing.
- * So P2 cannot enter.

Bounded waiting mutual exclusion with test-and-set.

```
do {
    waiting[i] = true;
    busy = true;
    while (!waiting[i] && !busy)
        busy = test-and-set (&lock);
```

waiting[i] = false;
/* critical section */

j = (i+1) % n;

while ((j != i) && !waiting[j])

j = (j+1) % n;

lock = F

No process in CS

lock = T

process in CS

$j \equiv i$
lock = false;

use

waiting [j] = false;
* remainder *

{ while (true);

Hardware solution - (Complete)

Date : _____

P1 P2 P3 P4

(case i) P1 wants to enter CS.

*

T	F	F	F
0	1	2	3

 key \equiv T

* while condition is satisfied, key execute TAS.

key = F, lock = T (CS is occupied)

* while not satisfied, exit.

*

F	F	F	F
---	---	---	---

, P1 enters in the CS.

(case ii) P2 wants to enter CS.

*

F	F	F	T
---	---	---	---

 key = T, lock = T

* while condition is satisfied, execute TAS

key = T, lock = T

* goes in infinite loop.

* if P1 completes execution,

* $j = (0+1) \% 4 = 1$, $i = 0$.

* ~~j = i~~ and keep incrementing j until $i = j$
or waiting [j] = T.

* since waiting [3] = T, waiting[3] = F (in if
wait)

* now, in P4, the while condition is not satisfied (since waiting[3] = F).

* so

F	F	F	F
---	---	---	---

 and P4 enters CS.

(case iii) P2 ~~wants~~ wants to enter CS.

*

F	T	F	F
---	---	---	---

 key = T lock = T

* while satisfied, execute TAS key = T lock = T.

* infinite loop.

* Repeat above case(ii)

X X

Initially, lock = F, key = F

waiting

F	F	F	F
0	1	2	3

.

- Note:
- If P4 is in CS, P3 no-tries to enter CS.
 - * once P4 completes execution, $j = (3+1) \% 4 = 0$, $i=3$.
 - * check while loop, satisfied, j increments
 $j = (1+1) \% 4 = 1$, $i=3$
 - * $j = 2$, $i=3$.
 $\hookrightarrow \text{waiting}[2] = T$.

Q) If P1 completes CS, and it wants to enter CS again, will P1 be allowed or denied?

Semaphore - SW tool for synchronization

- * integer value accessed/modified by wait() and signal().

```
wait(s){ // decrementing
    while(s <= 0), // busy wait
    s--;
}
```

```
CS
do{ // mutual exclusion
    wait(mutex); // block
    // CS
    signal(mutex); // continue
    // remainder.
} while(true);
```

```
signal(s){
    s++;
}
```

```
no busy waiting
* wait(semaphore *s){
    s->value--;
    if(s->value < 0){
        add process to s->list;
        block();
    }
}
```

```
Signal(semaphore *s){
    s->value++;
    if(s->value <= 0){
        remove a process P from
        s->list;
        wakeup(P);
    }
}
```

Bounded Buffer Problem

n buffers, each can hold one item

Semaphore mutex initialized to 1

Semaphore full initialized to 0

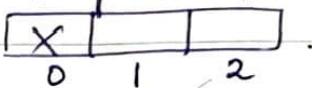
Semaphore empty initialized to n

<p><u>producer:</u></p> <pre>do{ /* produce an item in next-produced */ wait(empty); wait(mutex); /* add next produced to buffer */ signal(mutex); signal(full); } while(true);</pre>

- I (1) Initially $n = 3$: empty = 3, mutex = 1, full = 0
 (2) wait(empty) \rightarrow wait(3). (3) wait(mutex) \rightarrow wait(1)
~~empty = 2~~
 (4)

X	*	X
---	---	---

 (5) signal(0) mutex = 1 (6) signal(0)
 enter CS.
 \therefore mutex = 1, full = 1, empty = 2.



- II (1) wait(2) (2) wait(1) (3)

X	X	
---	---	--

 (4) signal(0)
~~empty = 1~~ mutex = 0
 (5) signal(1)
~~full = 2~~

case(i) When producer in CS, and one consumer tries to.

CS: mutex = 1/0 · + wait(full) \rightarrow wait(0)
 empty = 3/2 · full = -1 ·
 full = 0 · Add process to list
 block()

consumer:

```
do { wait(full);  

    wait(mutex);  

    /* remove item */  

    signal(mutex);  

    signal(empty);  

    /* consume item */  

} while(true);
```

now producer completes CS.
~~wait(mutex)~~

case ii) If the buffers are full & the producer tries to enter CS, is it allowed?

case iii) If buffers are empty & consumer tries to enter CS, is it allowed?

case iv) If consumer is in CS, the producer is called empty = 0
 is it allowed or denied? (producer full = 3)

case ii) In this case, empty mutex = 1, full = 3, empty = 0 ·
 * producers tries to execute code.

(2) wait(empty) \rightarrow wait(0) \therefore empty = -1 ·
 Adds this process to ~~list~~ list and
 blocks the process.

So producer will not be allowed

iii) In this case, mutex = 1, full = 0, empty = 3.

Consumer executes the code.

Date: _____

- wait(full) \rightarrow wait(0) \rightarrow full = -1.
adds process to list and blocks the process.

\therefore consumer is denied.

iv) mutex = 1, full = 3, empty = 0

consumer.

- wait(full) \rightarrow full = 2.
- wait(mutex) \rightarrow mutex = 0.

- CS

produces:

- wait(empty) \rightarrow empty = -1. (blocks the producer).

* Reader-writer problem

* Dining philosopher