

Docker & Containers

Vikram Krishnamurthy

Agenda

- *History – cgroups, namespaces, overlayfs*
- *What is/are Docker/Containers, what problems do they solve?*
- *How do Containers differ from Hypervisor based virtualization?*
- *What difference does Docker bring to Containers?*
- *How does Docker work fundamentally?*
- *Demo Time!!*

Resource Isolation


- ***Compute Resources:***
 - *CPU – Consume and Release*
 - *Memory – Consume and Release*
 - *N/W – Consume and Release*
 - *Disk/Secondary Storage – Consume and Hold*

Namespaces

- **Namespaces** are kernel enforced user space views.
- They are mechanisms to abstract, isolate, and limit the visibility that a group of processes has over various system entities such as process trees, network interfaces, user IDs, and filesystem mounts.
- Namespace Types:
 - Process ID – Allows a process to spin off a new process tree
 - Mount – confine set of mount points visible within a process namespace
 - UTS – isolating systems host and domain name
 - IPC – demarcate processes from using System V and POSIX message queries
 - User – allow a process to use unique UID and GID within and outside namespaces
 - Network – abstraction of network protocol services and interfaces
 - Cgroups – virtualizes content of /proc/self/cgroup file.

Cgroups


- Cgroups limit the amount of resources a process can consume (CPU, memory, network bandwidth, etc.). They were introduced in Linux kernel 2.6.24












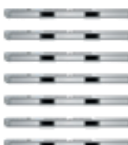



- Applying cgroups on namespaces allows isolation of processes into containers within a system, where resources can be managed distinctly.

OverlayFS

- Union file systems are a creative solution to allow a virtual merge of multiple folders, while keeping their actual contents separate. The Overlay file system (OverlayFS) is one example of these, though it is more of a mounting mechanism than a file system.
- OverlayFS allows you to overlay the contents (both files and directories) of one directory onto another
- *Lower* directory is Read only, *Upper* directory is RW



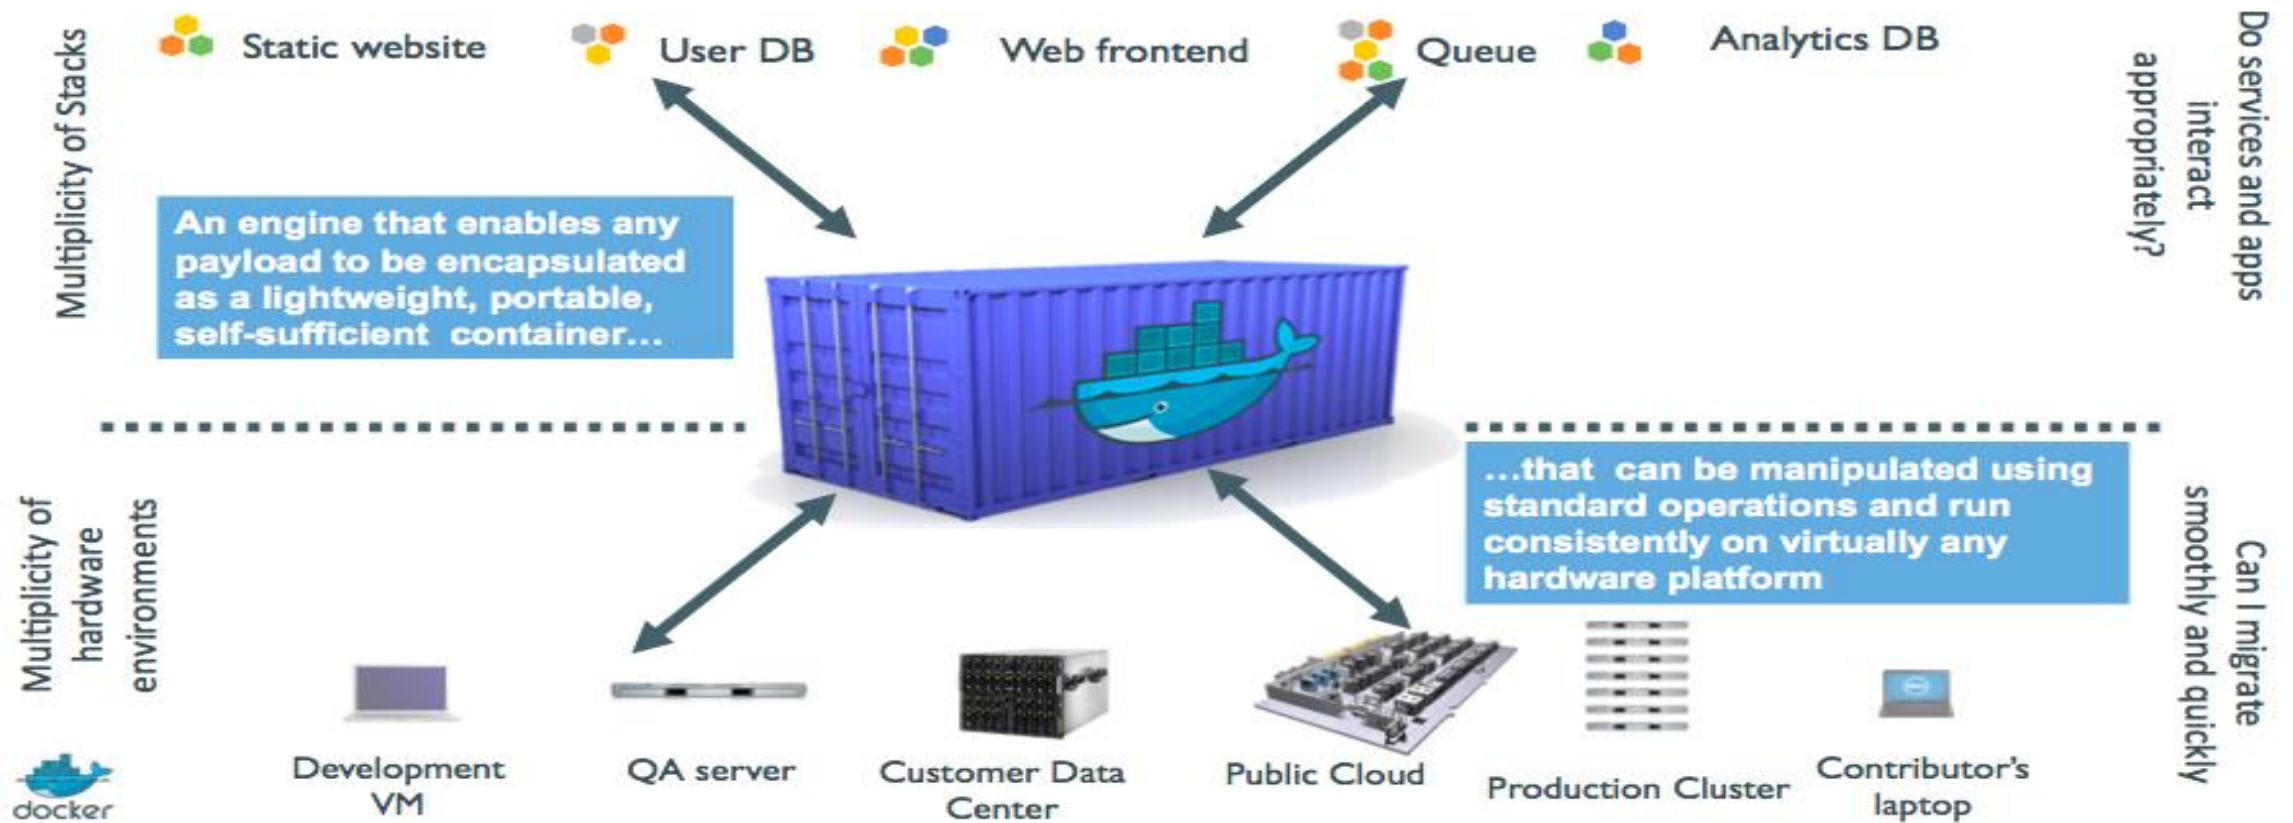
The Challenge – The Matrix From Hell

	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
								

Help From Elsewhere - Shipping Containers



Applied to IT World - Application Containers



Containers vs Processes vs VMs

Process

- Isolate address space
- No isolation for files or networks
- Lightweight

Virtual Machine

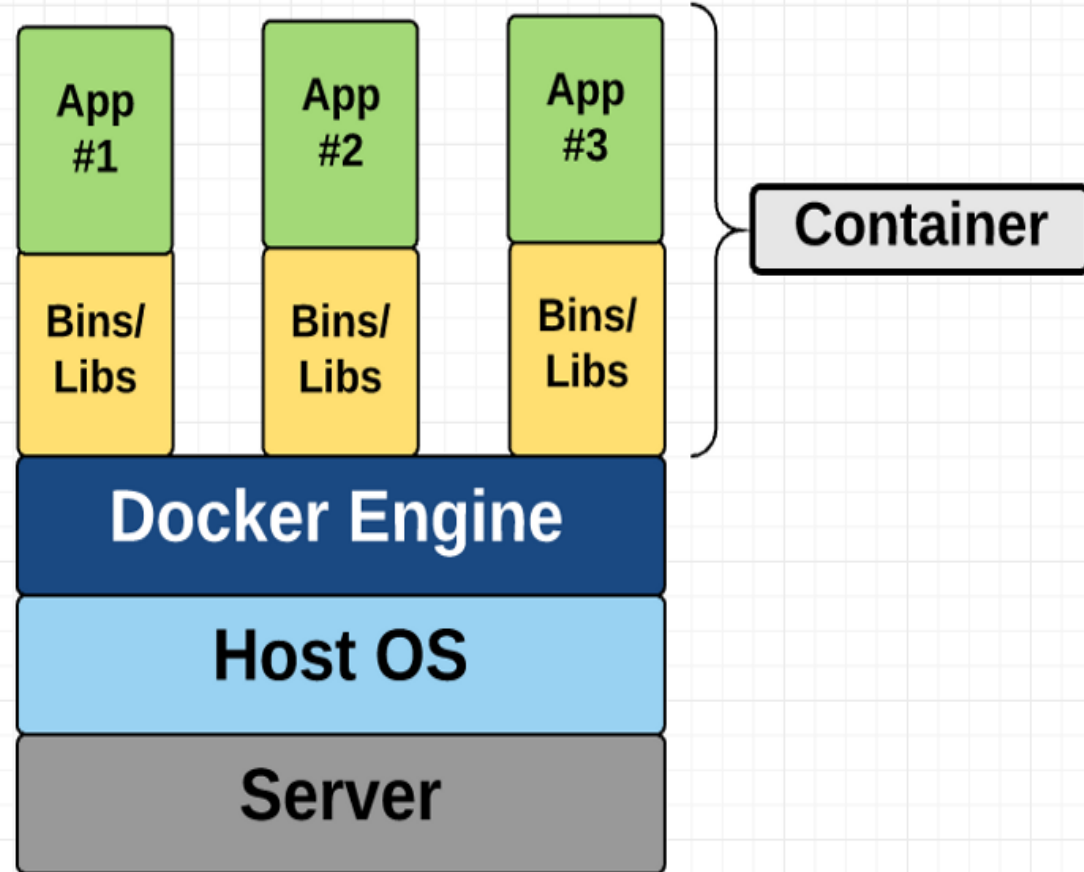
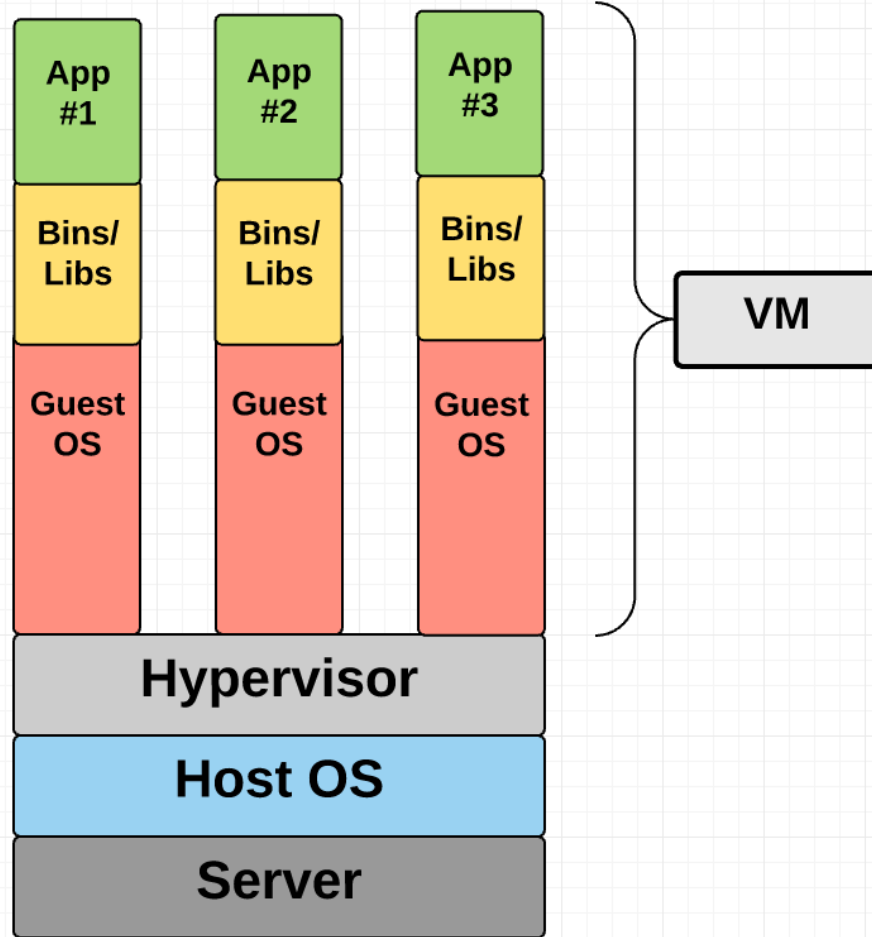
- Isolate address space
- isolate files and networks
- Heavyweight

Container

- Isolate address space
- isolate files and networks
- Lightweight

Containers vs Virtual Machines?

Illustrate with a diagram differences between container and virtual machine 6



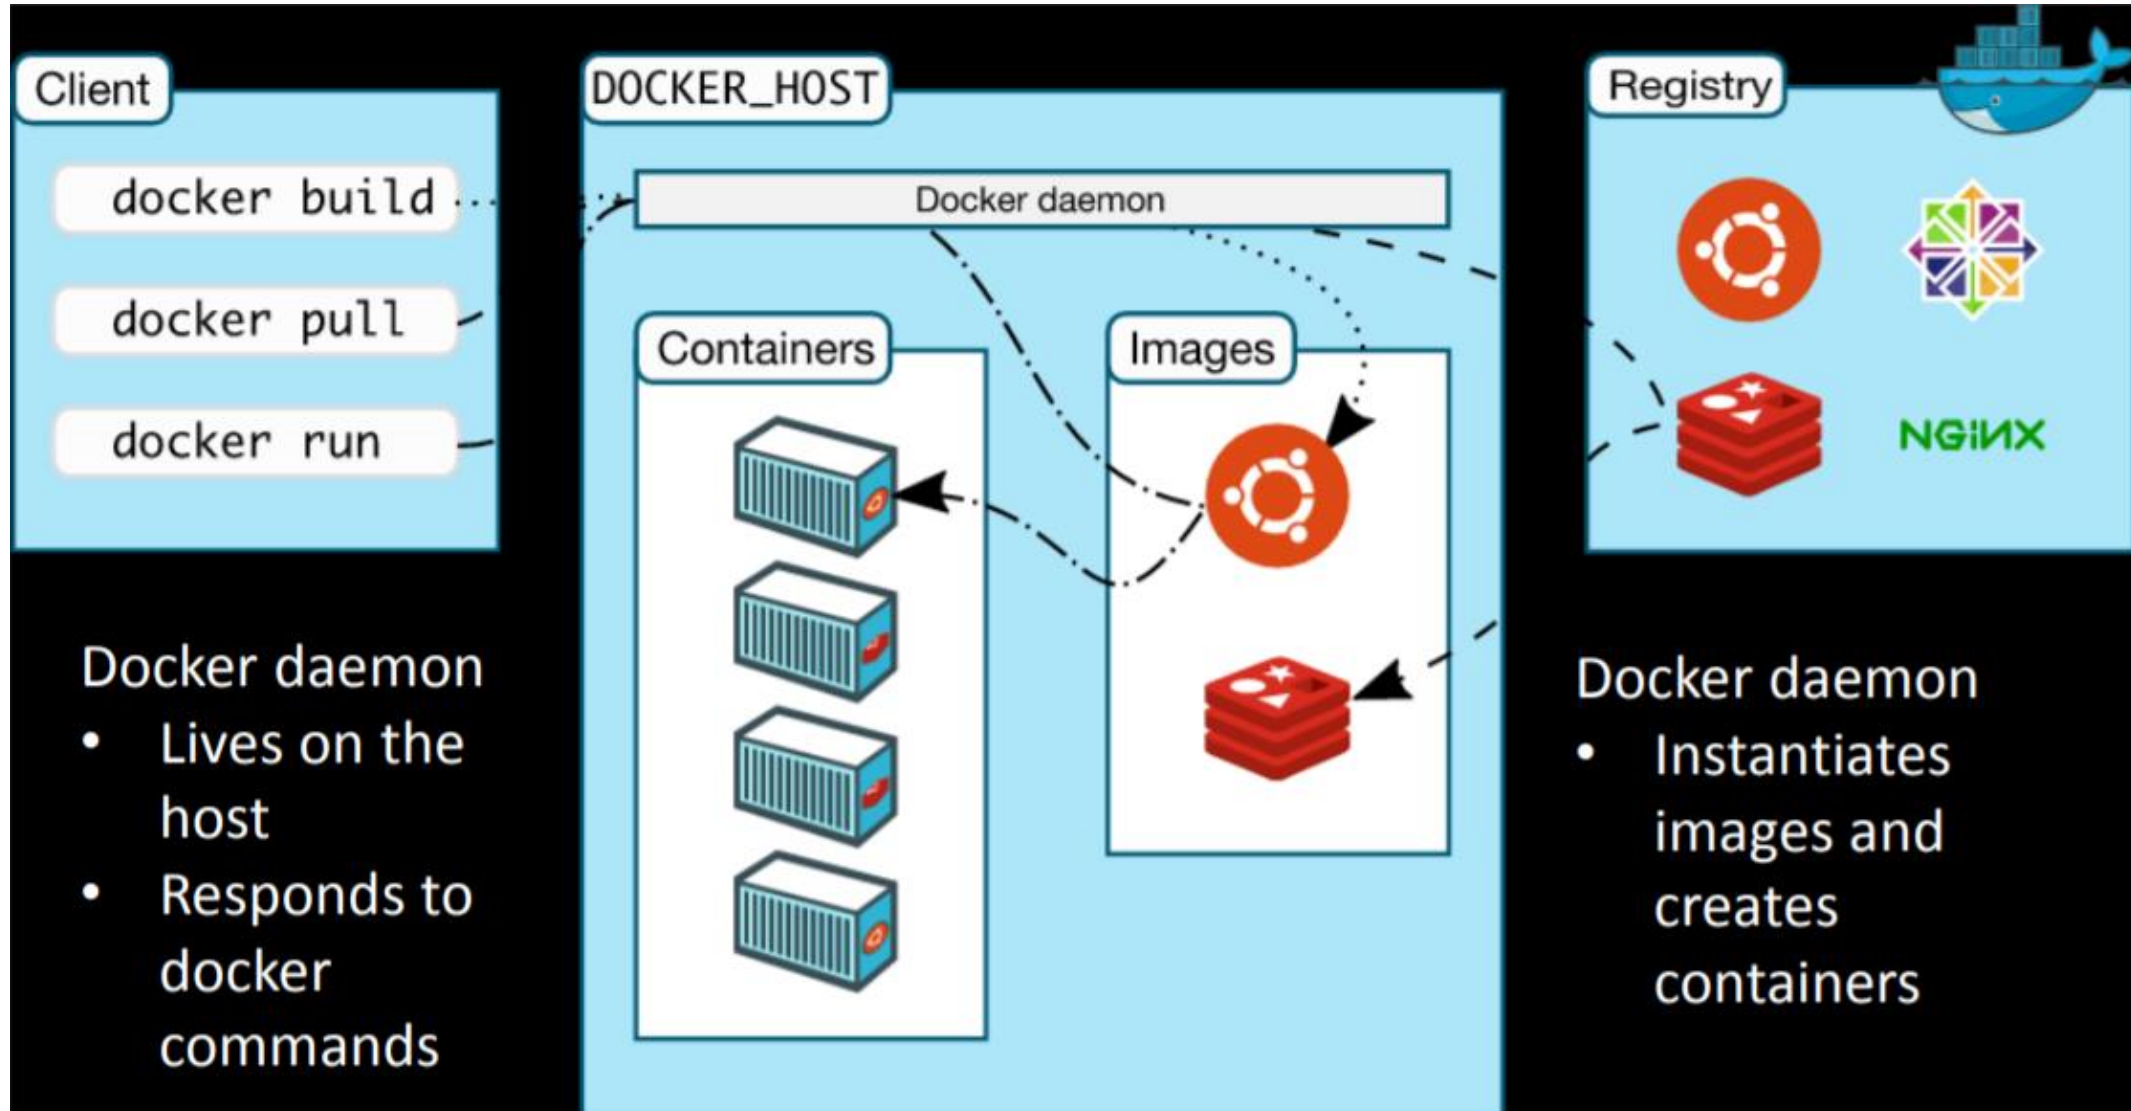
Hypervisors virtualize hardware, Containers virtualize OS!!

Why Should We Care about Containers?

- ❖ Build once, run *almost* anywhere
- ❖ A clean, safe, portable runtime environment for the app.
- ❖ Eliminate worries about missing dependencies, packages or compatibility between different platforms.
- ❖ Run each app in its own isolated container, so various versions of libraries and app can be run without conflicts.
- ❖ Automate testing, integration, packaging
- ❖ A VM without the overhead of a VM.

Docker Basic Architecture

With a neat diagram explain the architecture of dockers



Docker – Important Terminology

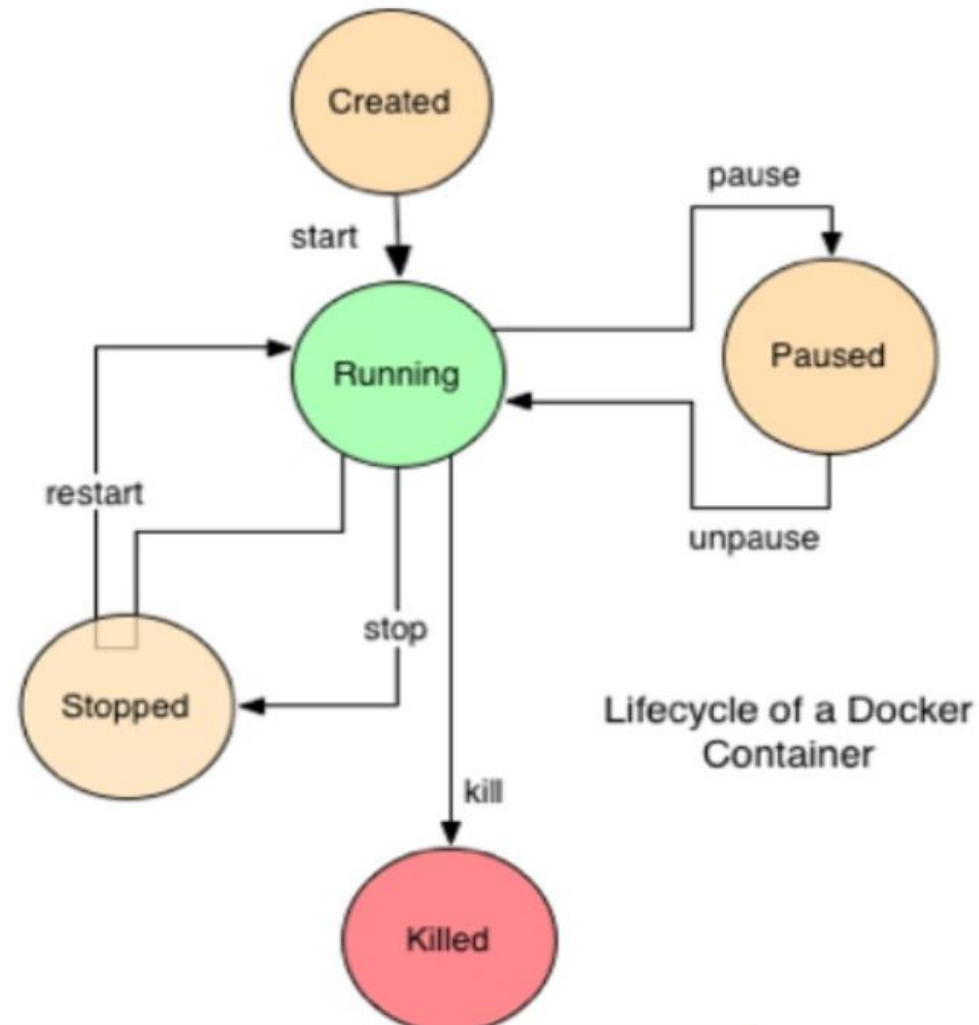
Define the following terms:

i) Image ii) Container iii) Dockerfile iv) Docker Client v) Docker Daemon/Engine

- Image: Persisted snapshot of application runtime and dependent binaries/libraries that can be «run»
- Container: Live running instance of a Docker «image»
- Dockerfile: A text document with commands to build Docker “image”.
- Docker Client : The utility that runs docker commands – *docker run*, *docker ps*, *docker build* etc.
- Docker Daemon/Engine: The server part that talks to the kernel, makes the system calls to create, operate and manage containers.

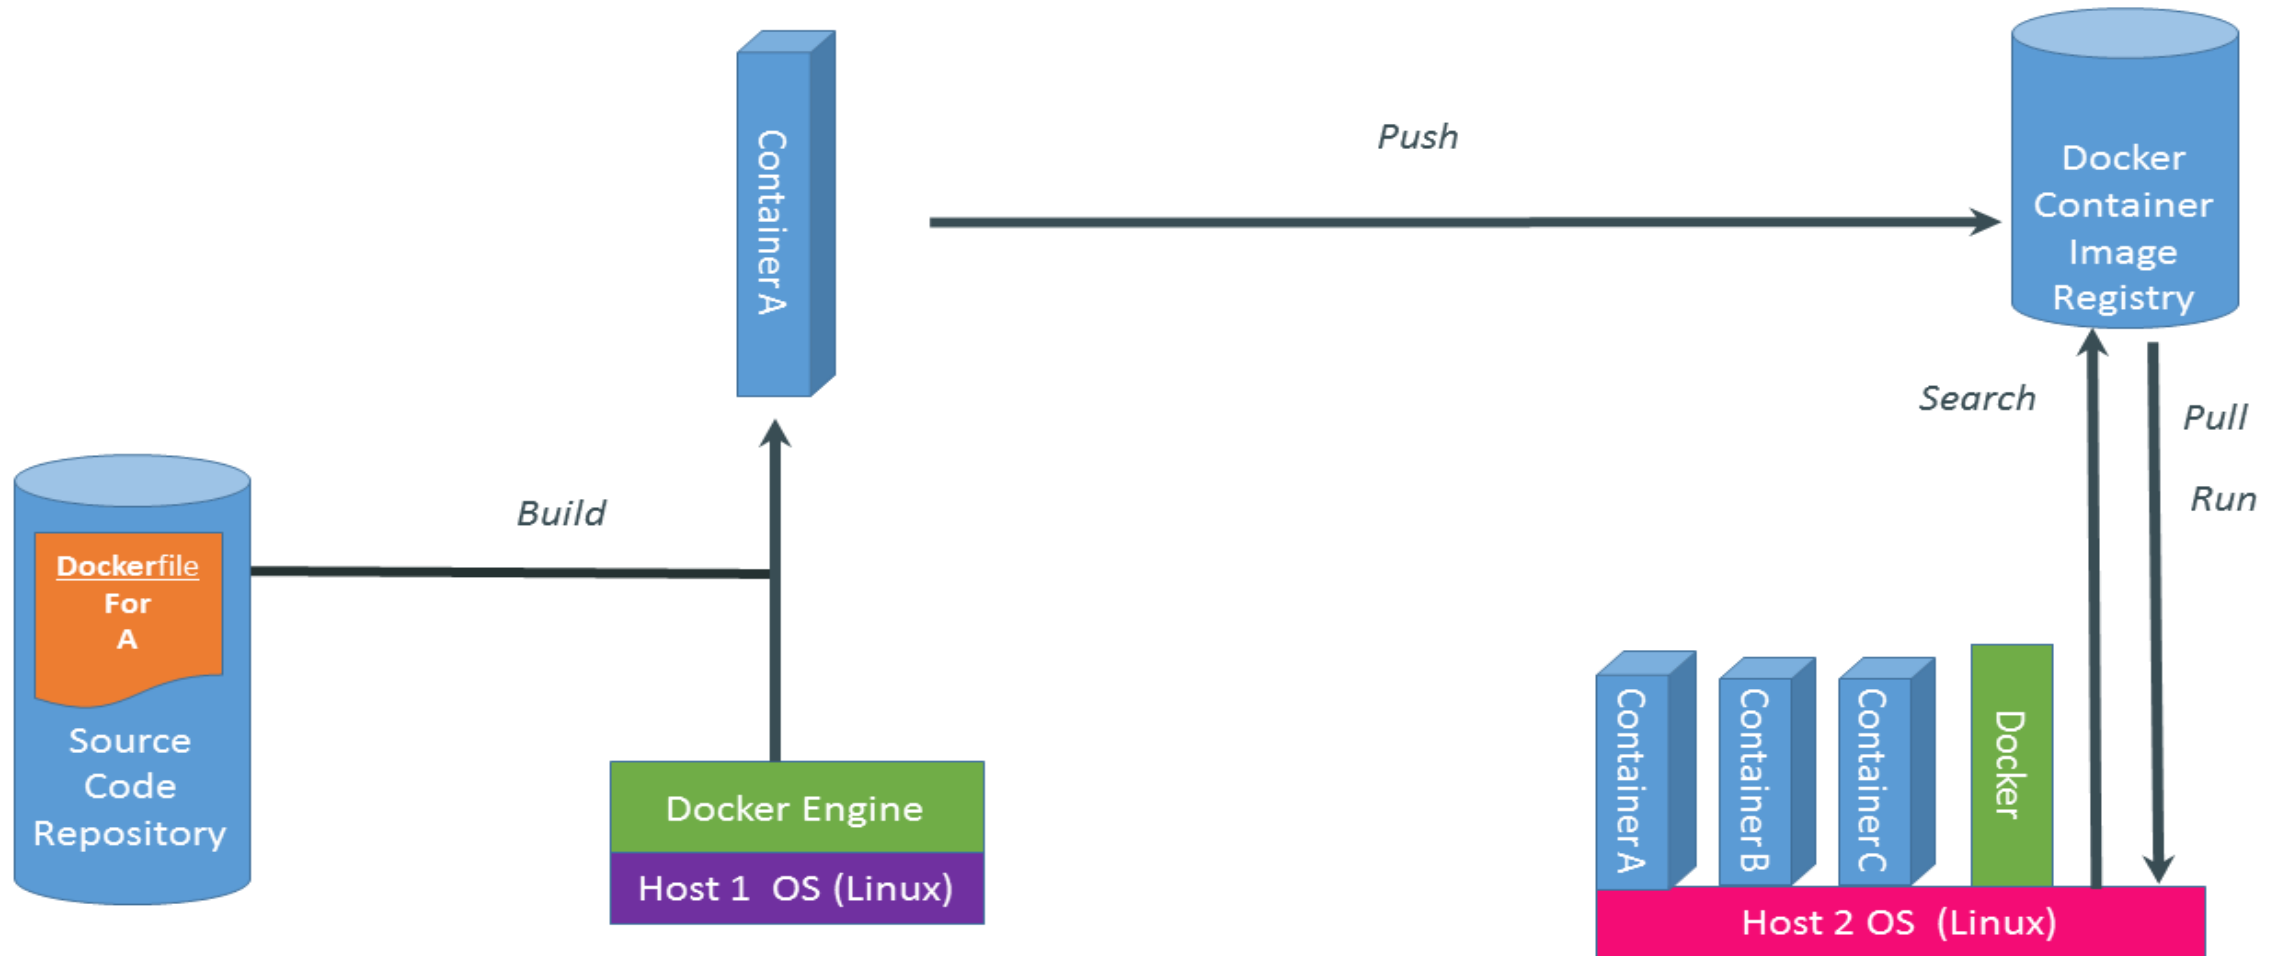
Docker Container Life Cycle

With a neat diagram describe Docker Container Life Cycle.



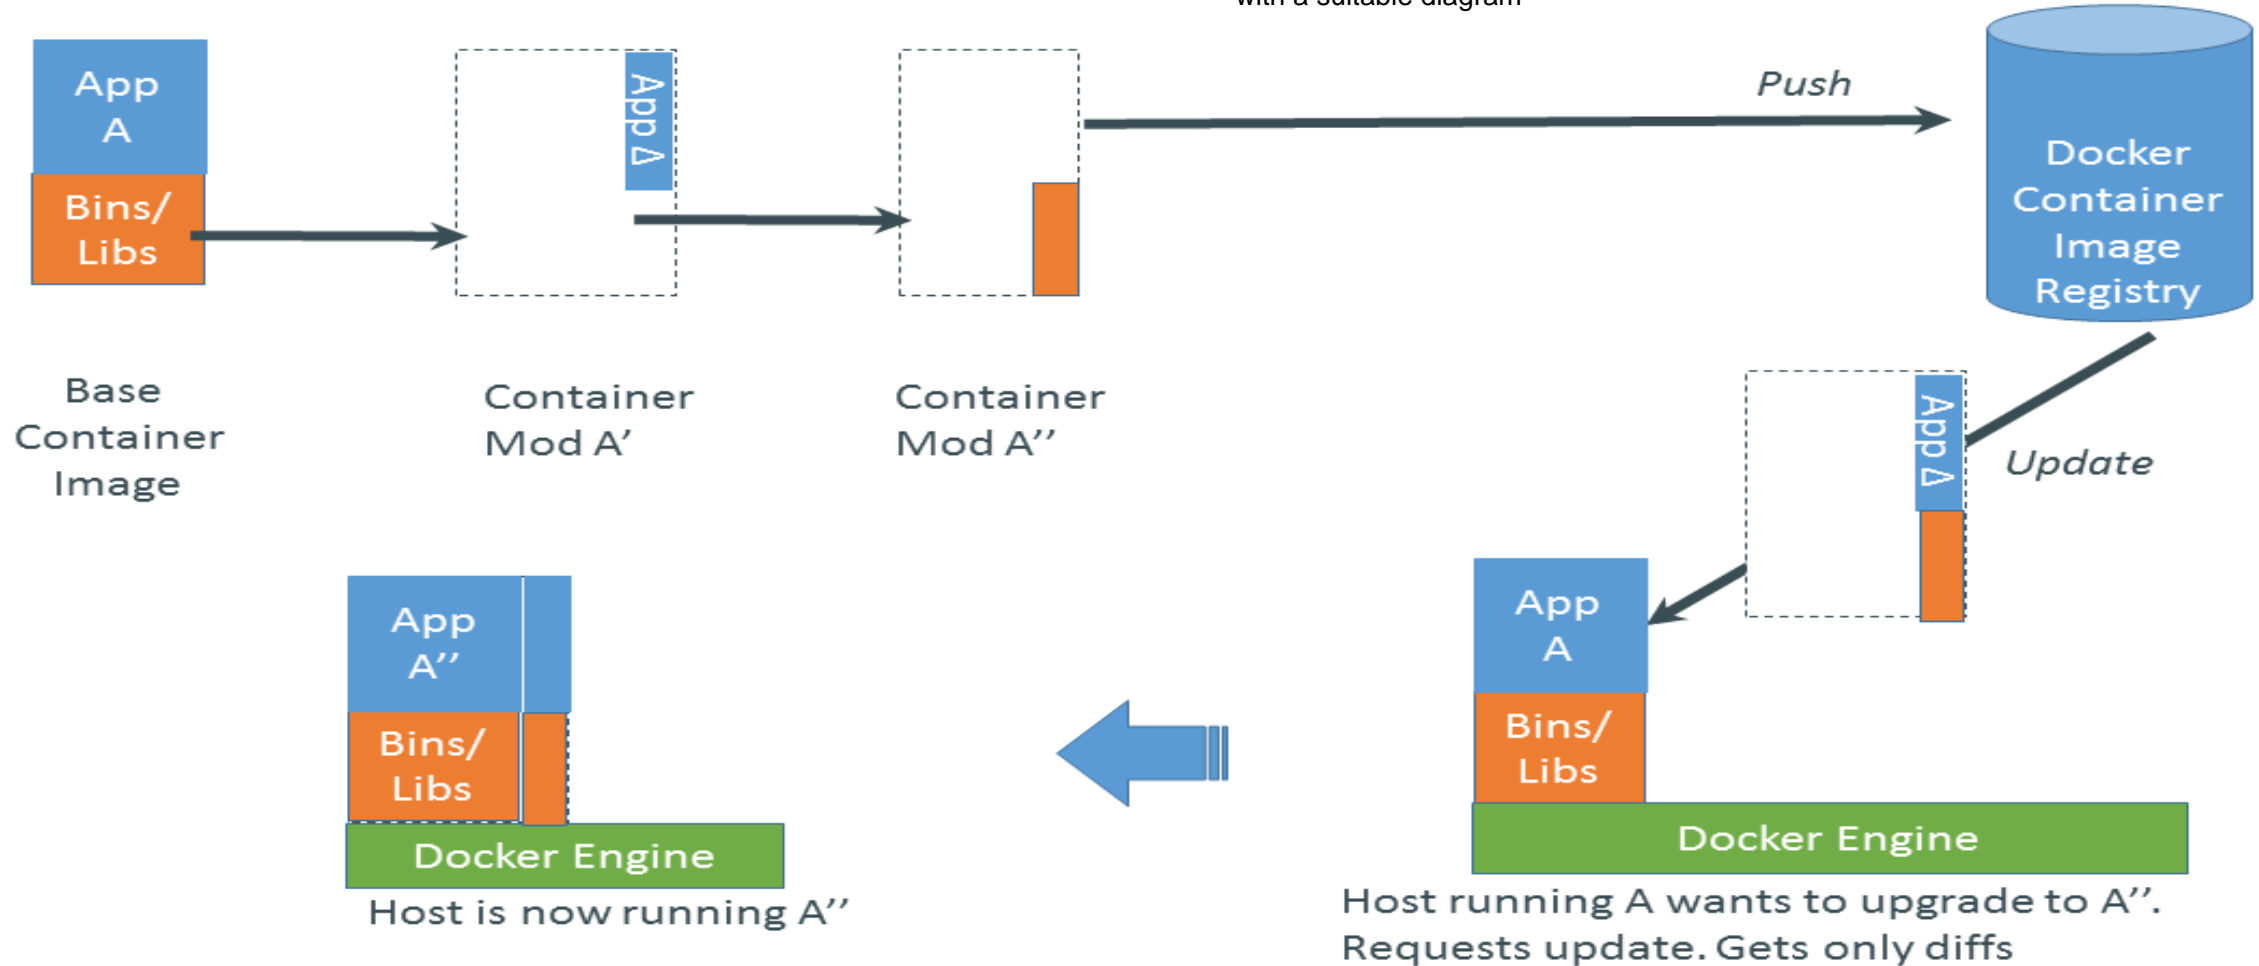
Docker Workflow - Basics

Illustrate basic workflow of the docker with a suitable diagram



Docker Workflow – App Updates / Changes

Illustrate workflow of the docker with App Updates / Changes with a suitable diagram



Docker – Examples

- Basics
 - `docker --version/version`
 - `docker pull`
 - `docker images`
- Container lifecycle
 - `docker run`
 - `docker ps`
 - `docker exec`
 - `docker pause/stop/kill/rm`
 - `docker unpause/restart`

Thank You.